

✓ Deep Learning Time Series Exam question

Problem Statement

In this Taks, you will explore a dataset containing the daily stock prices of a company, listed from December 1980 onwards. The data includes the following information for each day: opening price, highest price, lowest price, closing price, adjusted closing price, and the volume of shares traded. The primary focus will be on predicting future values of the Close price, using various time series analysis techniques.

✓ Step 1: Exploratory Data Analysis

- Print a few rows to understand the data's structure.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import LSTM, Dense
from sklearn.metrics import mean_squared_error, precision_score, recall_score, f1

import pandas as pd

df = pd.read_csv('TimeSeries.csv')
df.head()
```



	Date	Open	High	Low	Close	Adj Close	Volume
0	1980-12-12	0.128348	0.128906	0.128348	0.128348	0.099058	469033600
1	1980-12-15	0.122210	0.122210	0.121652	0.121652	0.093890	175884800
2	1980-12-16	0.113281	0.113281	0.112723	0.112723	0.086999	105728000
3	1980-12-17	0.115513	0.116071	0.115513	0.115513	0.089152	86441600
4	1980-12-18	0.118862	0.119420	0.118862	0.118862	0.091737	73449600




Next steps:

[Generate code with df](#)[View recommended plots](#)

- Plot the time series data to understand its characteristics (optional).

```
import matplotlib.pyplot as plt

plt.figure(figsize=(14, 7))
plt.plot(df['Date'], df['Close'])
plt.title('Time Series of Closing Prices')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.show()
```




- Observe the trends in the series.

```
df['Close'].plot(figsize=(10, 6))
plt.title('Close Price over Time')
plt.show()
```



Close Price over Time



✓ Step 2: Data Pre-processing and Feature Extraction

- Create the training and testing data.
- Make sure your data is in the right format before using it in your model.

Hint:

You may need to create a function to generate the training and testing datasets or use a library. Also, ensure you use the correct format and shape for the data.

Method example:

```
def prepare_data(df, feature_columns, target_column, sequence_length, test_size=0.2):  
    # Scale the features  
    scaler = MinMaxScaler(feature_range=(0, 1))  
    scaled_data = scaler.fit_transform(df[feature_columns])  
  
    # Create the sequences and targets  
    x, y = [], []  
    for i in range(sequence_length, len(df)):  
        x.append(scaled_data[i-sequence_length:i])
```

```
y.append(scaled_data[i, feature_columns.index(target_column)])

x, y = np.array(x), np.array(y)

# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=test_size, random

return X_train, X_test, y_train, y_test
```

Note: This method is provided as an example and may not be directly applicable to your specific model.

- Check the columns if any of them needs to be converted to another data type.

```
df['Date'] = pd.to_datetime(df['Date'])
df.set_index('Date', inplace=True)
```



Show hidden output

Next steps: [Explain error](#)

```
df.head()
```



	Open	High	Low	Close	Adj Close	Volume	
Date							
1980-12-12	0.128348	0.128906	0.128348	0.128348	0.099058	469033600	
1980-12-15	0.122210	0.122210	0.121652	0.121652	0.093890	175884800	
1980-12-16	0.113281	0.113281	0.112723	0.112723	0.086999	105728000	
1980-12-17	0.115513	0.116071	0.115513	0.115513	0.089152	86441600	
1980-12-18	0.118862	0.119420	0.118862	0.118862	0.091737	73449600	

Next steps: [Generate code with df](#) [View recommended plots](#)

- Check and handle missing values if there is any.

```
df.isnull().sum()
df.fillna(method='ffill', inplace=True)
```

```
df.isnull().sum()
```

```

⇒ Open      0
   High     0
   Low      0
   Close    0
   Adj Close 0
   Volume   0
dtype: int64

```

- Perform any additional necessary preprocessing steps including normalization.

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(df[['Close']])
```

- Create the training and testing data.

```

# def prepare_data(df, feature_columns, target_column, sequence_length, test_size=
#     scaler = MinMaxScaler(feature_range=(0, 1))
#     scaled_data = scaler.fit_transform(df[feature_columns])

#     x, y = [], []
#     for i in range(sequence_length, len(df)):
#         x.append(scaled_data[i-sequence_length:i])
#         y.append(scaled_data[i, feature_columns.index(target_column)])

#     x, y = np.array(x), np.array(y)
#     X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=test_siz
#     return X_train, X_test, y_train, y_test

def prepare_data(df, feature_columns, target_column, sequence_length, test_size=0.
    scaled_data = scaler.fit_transform(df[feature_columns])
    x, y = [], []
    for i in range(sequence_length, len(df)):
        x.append(scaled_data[i-sequence_length:i])
        y.append(scaled_data[i, feature_columns.index(target_column)])

    x, y = np.array(x), np.array(y)
    X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=test_size,
    return X_train, X_test, y_train, y_test

```

Double-click (or enter) to edit

```
X_train, X_test, y_train, y_test = prepare_data(df, ['Close'], 'Close', sequence_
```

✓ Step 4: Model Design

- Define and select the appropriate time series models (e.g., RNN, LSTM, etc.).
- Use appropriate activation functions for the hidden layers and the output layer.

```
from keras.models import Sequential
from keras.layers import LSTM, Dense
```

```
model = Sequential()
model.add(LSTM(units=50, return_sequences=True, input_shape=(X_train.shape[1], X_
model.add(LSTM(units=50, return_sequences=False))
model.add(Dense(units=1))
```

✓ Step 5: Model Compilation

- Choose a suitable optimizer. Explain your choice.
- Select an appropriate loss function. Explain why this loss function is optimal for your model.
- Compile your ANN model.

```
model.compile(optimizer='adam', loss='mean_squared_error')
```

✓ Step 6: Model Training

- Train your model on the training data. Set appropriate values for the number of epochs and batch size.
- Ensure to include validation data.
- Evaluate your model's performance during training.

```
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=
```

```
⇒ Epoch 1/10
218/218 [=====] - 15s 67ms/step - loss: 5.8613e-05 -
Epoch 2/10
218/218 [=====] - 13s 58ms/step - loss: 6.3334e-05 -
Epoch 3/10
218/218 [=====] - 13s 58ms/step - loss: 6.2226e-05 -
Epoch 4/10
218/218 [=====] - 13s 58ms/step - loss: 5.8026e-05 -
Epoch 5/10
218/218 [=====] - 12s 57ms/step - loss: 6.0653e-05 -
Epoch 6/10
218/218 [=====] - 12s 57ms/step - loss: 5.5068e-05 -
Epoch 7/10
218/218 [=====] - 12s 54ms/step - loss: 4.8920e-05 -
```

```
Epoch 8/10
218/218 [=====] - 14s 62ms/step - loss: 7.4077e-05 -
Epoch 9/10
218/218 [=====] - 12s 54ms/step - loss: 4.9272e-05 -
Epoch 10/10
218/218 [=====] - 12s 55ms/step - loss: 4.4391e-05 -
```

✓ Step 7: Model Evaluation

- Use the test set and cross-validation to evaluate your model's performance comprehensively.
- Show the performance metrics's results (Precision, Recall, and F1 Score).
- Discuss the results and any potential areas for improvement.
- Use the test set and cross-validation to evaluate your model's performance comprehensively.

```
predictions = model.predict(X_test)
mse = mean_squared_error(y_test, predictions)
print('Mean Squared Error:', mse)
```

```
⇒ 69/69 [=====] - 2s 25ms/step
Mean Squared Error: 4.013534857333456e-05
```

- Show the performance metrics's results (Precision, Recall, and F1 Score).

```
print(f'Type of y_test: {type(y_test[0])}')
print(f'Type of predictions: {type(predictions[0])}')
```

```
⇒ Type of y_test: <class 'numpy.float64'>
Type of predictions: <class 'numpy.ndarray'>
```

```
binary_predictions = (predictions > 0.5).astype(int)
```

```
if y_test.dtype == 'float64' and len(set(y_test)) > 2:
    binary_y_test = (y_test > 0.5).astype(int)
else:
    binary_y_test = y_test
```

```
binary_predictions = (predictions > 0.5).astype(int)
```

```
from sklearn.metrics import precision_score, recall_score, f1_score

precision = precision_score(binary_y_test, binary_predictions)
recall = recall_score(binary_y_test, binary_predictions)
f1 = f1_score(binary_y_test, binary_predictions)

print(f'Precision: {precision}, Recall: {recall}, F1 Score: {f1}')
```

⇒ Precision: 1.0, Recall: 1.0, F1 Score: 1.0

- Discuss the results and any potential areas for improvement.

```
print("The model has a precision of", precision, ", recall of", recall, "and F1 score of", f1)
print("Potential areas for improvement could include adjusting the model architecture")
```

⇒ The model has a precision of 1.0 , recall of 1.0 and F1 score of 1.0
Potential areas for improvement could include adjusting the model architecture

Start coding or [generate](#) with AI.