

# Machine Learning (classification problem) Exam question

## Problem Statement

In this task, you are provided with a dataset that includes several features,  
reviews this column has keywords of the review

rating this column has reviews of the app

Source this column has store the app was downloaded from

Country\_code this column indicate to which country the reviewer is from

Your task is to clean, visualize and apply a machine learning model to accurately predict this outcome and evaluate the performance of your model. Your targeted column is the review column.

## ✓ Step 1: Import necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

## ✓ Step 2: Read the Dataset

```
# Load the dataset
df = pd.read_csv('Exam_reviews.csv')
df.head()
```



reviews rating

source country\_code



0 ['cluttered', 'confusing', 'put', 'back'] 3

1 ['blood', 'work', 'impossible', 'understand', ...] 2

2 ['love', 'mychart', 'yrs', 'access', 'records'...] 5

3 ['enjoy', 'using', 'mychart', 'easy', 'use', '...'...] 5

country\_code

الرقم الدولي



Next steps:

[Generate code with df](#)[View recommended plots](#)

## ✓ Step 3: Explore the dataset

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4689 entries, 0 to 4688
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   reviews         4689 non-null   object
1   rating           4689 non-null   int64
2   source           4689 non-null   object
3   country_code     4689 non-null   object
dtypes: int64(1), object(3)
memory usage: 146.7+ KB
```


# Check for missing values

`df.isnull().sum()`

```
reviews      0
rating       0
source       0
country_code 0
dtype: int64
```

`df.describe()`



	rating	
count	4689.000000	
mean	3.843890	
std	1.618982	
min	1.000000	
25%	3.000000	
50%	5.000000	
75%	5.000000	
max	5.000000	

## ✓ Step 4: preprocess the dataset (if needed)


```
# Drop any missing values
df.dropna(inplace=True)

# Example preprocessing step for text data
vectorizer = CountVectorizer(stop_words='english')
X = vectorizer.fit_transform(df['reviews'])


# Define the target variable
y = df['rating']
```

## ✓ Step 5: Visualize and see the relations between the features


```
# Visualize the distribution of ratings
plt.figure(figsize=(10, 6))
sns.countplot(x='rating', data=df)
plt.title('Distribution of Ratings')
plt.show()
```



```
# Visualize the distribution of reviews by source
plt.figure(figsize=(10, 6))
sns.countplot(x='source', data=df)
plt.title('Distribution of Reviews by Source')
plt.show()
```



```
# Visualize the distribution of reviews by country
plt.figure(figsize=(10, 6))
sns.countplot(x='country_code', data=df)
plt.title('Distribution of Reviews by Country')
plt.show()
```



## ✓ Step 6: split the dataset and start training

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st

# Train the model
model = MultinomialNB()
model.fit(X_train, y_train)
```



▼ MultinomialNB  
MultinomialNB()

## ✓ Step 7: use cross-validation and Test your model

```
# Cross-validation
cv_scores = cross_val_score(model, X_train, y_train, cv=5)
print("Cross-validation scores:", cv_scores)
print("Mean cross-validation score:", np.mean(cv_scores))
```



Cross-validation scores: [0.74567244 0.72133333 0.71466667 0.73066667 0.73333  
Mean cross-validation score: 0.7291344873501997

```
# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Classification report
print("Classification Report:\n", classification_report(y_test, y_pred))
```